



Memory Corruption Vulnerabilities, Part I



Some Terminology

- Software **error**
 - A programming mistake that make the software not meet its expectation
- Software **vulnerability**
 - A software error that can lead to possible attacks
- **Attack**
 - The process of exploiting a vulnerability
 - An attack can exploit a vulnerability to achieve additional functionalities for attackers
 - E.g., privilege escalation, arbitrary code execution



Software, One of the Weakest Links in the Security Chain

- Cryptographic algorithms are strong
 - Nobody attacks it
 - Even for crypto hash
- However, even for the best crypto algorithms
 - Software has to implement them correctly
 - A huge of amount of software for other purposes
 - Access control; authentication; ...
- Which programming language to use also matters



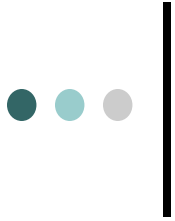
Language of Choice for System Programming: C/C++

- Systems software
 - OS; hypervisor; web servers; firmware; network controllers; device drivers; compilers; ...
- Benefits of C/C++: programming model close to the machine model; flexible; efficient
- BUT **error-prone**
 - Debugging memory errors is a headache
 - Perhaps on par with debugging multithreaded programs
 - Huge security risk



Agenda

- Compare C to Java
- Common errors for handling C-style buffers
- How to exploit buffer overflows: stack smashing



Comparing C to Java: language matters for security



Comparing C to Java

- Their syntax very similar
- Type safety
 - Safety: something “bad” won’t happen
 - “No untrapped errors”
- Java is type safe
 - Static type system + runtime checks + garbage collection
- C is type unsafe
 - Out-of-bound array accesses
 - Manual memory management
 - Bad type casts
 - ...



Java: Runtime Array Bounds Checking

- Example:

```
int a[10];
```

```
a[10] = 3;
```


- An exception is raised
- The length of the array is stored at runtime (the length never changes)

Java: Runtime Array Bounds Checking

- Java optimizer can optimize away lots of unnecessary array bounds checks

```
int sum = 0;  
for (i = 0; i < a.length; i++) {  
    sum += a[i];  
}
```

bounds checking unnecessary



C: No Array Bounds Checking

```
int a[10];
```

```
a[10] = 3;
```

- Result in a silent error in C (buffer overflow)
- After that, anything can happen
 - Mysterious crash depending on what was overwritten
 - A security risk as well: if the data written can be controlled by an attacker, then he can possibly exploit this for an attack



Memory Management

- C: manual memory management
 - malloc/free
 - Memory mismanagement problems: use after free; memory leak; double frees
- Java: Garbage Collection
 - No “free” operations for programmers
 - GC collects memory of objects that are no longer used
 - Java has no problems such as use after free, as long as the GC is correct



Non-Null Checking and Initialization Checking in Java

- An object reference is either valid or null
 - Automatic non-null checking whenever it's used
 - once again, optimizers can eliminate many non-null checks
 - Example: `A a = new A(); a.f = 3;`
- A variable is always initialized before used
 - Java has a static verifier (at the bytecode level) that guarantees this




Java Strings

- Similar to an array of chars, but immutable
 - The length of the string is stored at runtime to perform bounds checking
- All string operations do not modify the original string (a la functional programming)
 - E.g., `s.toLowerCase()` returns a new string

● ● ● | C-Style Strings

- C-style strings consist of a contiguous sequence of characters, terminated by and including the first null character.
 - String length is the number of bytes preceding the null character.
 - The number of bytes required to store a string is the number of characters plus one (times the size of each character).

h	e	l	l	o	\0
---	---	---	---	---	----



C Strings: Usage and Pitfalls

● ● ● | Using Strings in C

- C provides many string functions in its libraries (libc)
- For example, we use the strcpy function to copy one string to another:

```
#include <string.h>
char string1[] = "Hello, world!";
char string2[20];
strcpy(string2, string1);
```




Using Strings in C

- Another lets us compare strings

```
char string3[] = "this is";  
char string4[] = "a test";  
if(strcmp(string3, string4) == 0)  
    printf("strings are equal\n");  
else printf("strings are different\n")
```

- This code fragment will print "strings are different". Notice that strcmp does **not** return a boolean result.

●●● | Other Common String Functions

- strlen: getting the length of a string
- strncpy: copying with a bound
- strcat/strncat: string concatenation
- gets, fgets: receive input to a string
- ...

Common String Manipulation Errors

- Programming with C-style strings, in C or C++, is error prone
- Common errors include
 - Buffer overflows
 - null-termination errors
 - off-by-one errors
 - ...

● ● ● | gets: Unbounded String Copies

- Occur when data is copied from an unbounded source to a fixed-length character array

```
void main(void) {  
    char Password[8];  
    puts("Enter a 8-character password:");  
    gets(Password);  
    printf("Password=%s\n", Password);  
}
```

● ● ● | strcpy and strcat

- The standard string library functions do not know the size of the destination buffer

```
int main(int argc, char *argv[]) {  
    char name[2048];  
    strcpy(name, argv[1]);  
    strcat(name, " = ");  
    strcat(name, argv[2]);  
    ...  
}
```

●●● | Better String Library Functions

- Functions that restrict the number of bytes are often recommended
- Never use `gets(buf)`
 - Use `fgets(buf, size, stdin)` instead

From gets to fgets


- `char *fgets(char *BUF, int N, FILE *FP);`
 - *“Reads at most N-1 characters from FP until a newline is found. The characters including to the newline are stored in BUF. The buffer is terminated with a 0.”*

```
void main(void) {  
    char Password[8];  
    puts("Enter a 8-character password:");  
    fgets(Password, 8, stdin);  
    ...  
}
```



Better String Library Functions

- Instead of `strcpy()`, use `strncpy()`
- Instead of `strcat()`, use `strncat()`
- Instead of `sprintf()`, use `snprintf()`



But Still Need Care

- `char *strncpy(char *s1, const char *s2, size_t n);`
 - *“Copy not more than n characters (including the null character) from the array pointed to by $s2$ to the array pointed to by $s1$; If the string pointed to by $s2$ is shorter than n characters, null characters are appended to the destination array until a total of n characters have been written.”*
 - What happens if the size of $s2$ is n or greater
 - It gets truncated
 - **And $s1$ may not be null-terminated!**

● ● ● | Null-Termination Errors

```
int main(int argc, char* argv[]) {  
    char a[16], b[16];  
    strncpy(a, "0123456789abcdef", sizeof(a));  
    printf("%s\n",a);  
    strcpy(b, a);  
}
```

a[] not properly terminated. Possible segmentation fault
if printf("%s\n",a);

How to fix it?

● ● ● | strcpy to strncpy

- Don't replace

`strcpy(dest, src)`

by

`strncpy(dest, src, sizeof(dest))`

but by

`strncpy(dest, src, sizeof(dest)-1)`

`dst[sizeof(dest)-1] = '\0';`

if dest should be null-terminated!

- You never have this headache in Java



Signed vs Unsigned Numbers

```
char buf[N];  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len > N)  
    {error ("invalid length"); return; }  
read(fd, buf, len);
```

We forget to check for negative lengths

len cast to unsigned and negative length overflows



Checking for Negative Lengths


```
char buf[N];  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len > N || len < 0)  
    {error ("invalid length"); return; }  
read(fd, buf, len);
```

It still has a problem
if the buf is going to be treated as a C string.



A Good Version

```
char buf[N];  
int i, len;  
  
read(fd, &len, sizeof(len));  
if (len > N-1 || len < 0)  
    {error ("invalid length"); return; }  
read(fd, buf, len);  
buf[len] = '\0'; // null terminate buf
```



Buffer Overflows



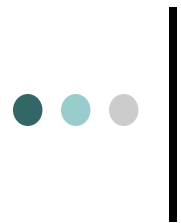
Problems Caused by Buffer Overflows

- The first Internet worm, and many subsequent ones (CodeRed, Blaster, ...), exploited buffer overflows
- Buffer overflows cause in the order of 50% of all security alerts
 - E.g., check out CERT, cve.mitre.org, or bugtraq
- Trends
 - Attacks are getting cleverer
 - defeating ever more clever countermeasures
 - Attacks are getting easier to do, by script kiddies




How Can Buffer Overflow Errors Lead to Software Vulnerabilities?

- All the examples look like simple *programming bugs*
- How can they possibly enable attackers to do bad things?
 - **Stack smashing** to exploit buffer overflows
 - Illustrate the technique using the Intel x86-64 architecture




Compilation, Program, and Process

- Compilation
 - From high-level programs to low-level machine code
- Program: static code and data
- Process: a run of a program



Process Memory Region



- Text: static code
- Data: also called heap
 - static variables
 - dynamically allocated data (malloc, new)
- Stack: program execution stacks



Program Stack

- For implementing procedure calls and returns
- Keep track of program execution and state by storing
 - local variables
 - Some arguments to the called procedure (callee)
 - Depending on the calling convention
 - return address of the calling procedure (caller)
 - ...

Stack Segment

The stack supports
nested invocation calls

Information pushed on
the stack as a result of
a function call is called
a frame

```
    b() {...}
    a() {
        b();
    }
    main() {
        a();
    }
```

Low memory

Unallocated

Stack frame
for b()

Stack frame
for a()

Stack frame
for main()

High memory

A stack frame is
created for each
subroutine and
destroyed upon
return.



Stack Frames

- Stack grows from high mem to low mem
- The stack pointer points to the top of the stack
 - RSP in Intel x86-64
- The frame pointer points to the end of the current frame
 - also called the base pointer
 - RBP in Intel x86-64
- The stack is modified during
 - function calls
 - function initialization
 - returning from a function



A Running Example

```
void function(int a, int b) {  
    char buffer[12];  
    gets(buffer);  
    return;  
}
```

```
void main() {  
    int x;  
    x = 0;  
    function(1,2);  
    x = 1;  
    printf("%d\n",x);  
}
```

Run “gcc -S -o example.s example.c”
to see its assembly code

- The exact assembly code will depend on many factors (the target architecture, optimization levels, compiler options, etc);
- We show the case for unoptimized x86-64

Function Calls

function (1,2)


movl	\$2, %esi
movl	\$1, %edi
call	function

pass the 2nd arg

pass the 1st arg


push the ret addr onto the stack,
and jumps to the function

Note: in x86-64, the first 6 args are passed
via registers (rdi, rsi, rdx, rcx, r8, r9)




Function Calls: Stacks

Before



After



Function Initialization

```
void function(int a, int b) {
```


<code>pushq %rbp</code>
<code>movq %rsp, %rbp</code>
<code>subq \$32, %rsp</code>

save the frame pointer

set the new frame pointer


allocate space for local
variables

Procedure prologue

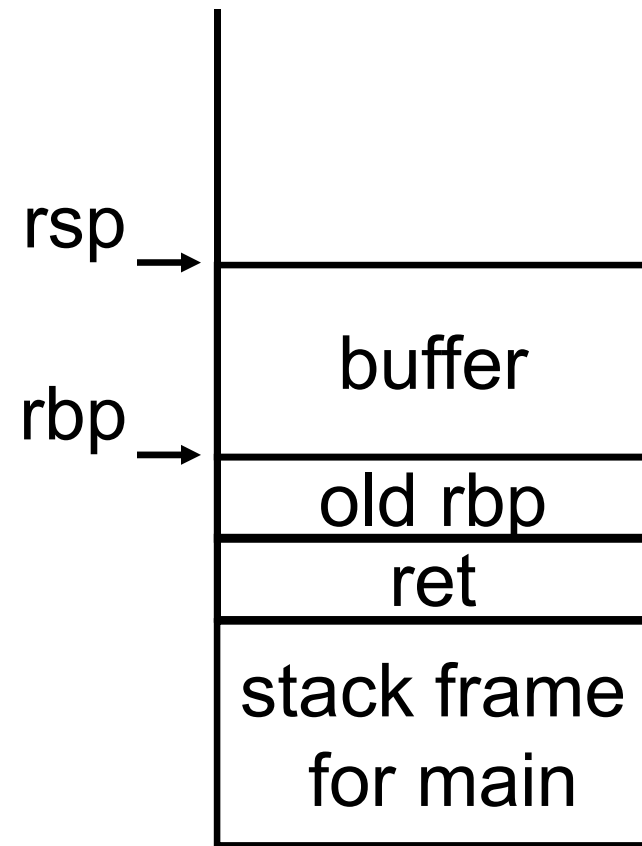


Function Initialization: Stacks

Before



After





Function Return


return;

movq %rbp, %rsp
popq %rbp
ret

restores the old stack pointer

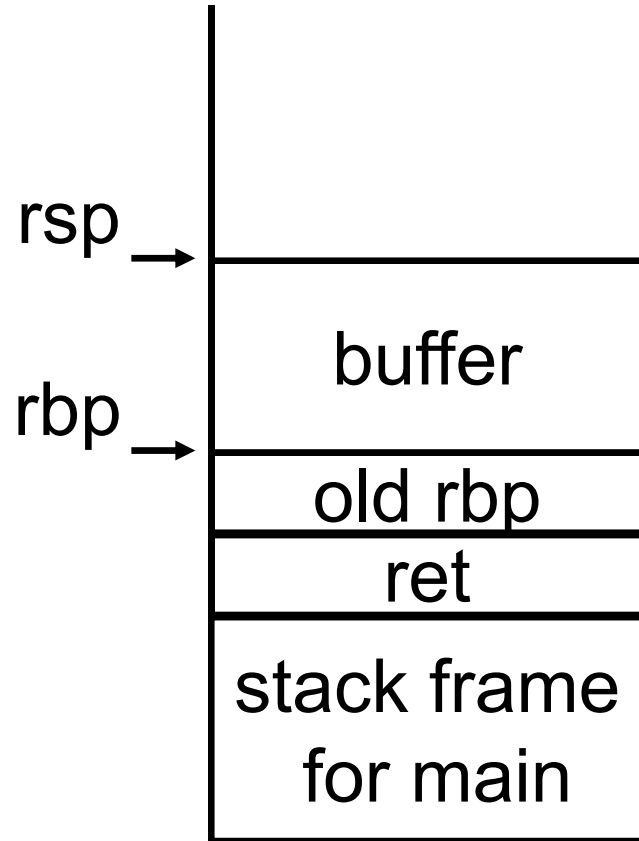
restores the old frame pointer

gets the return address, and
jumps to it

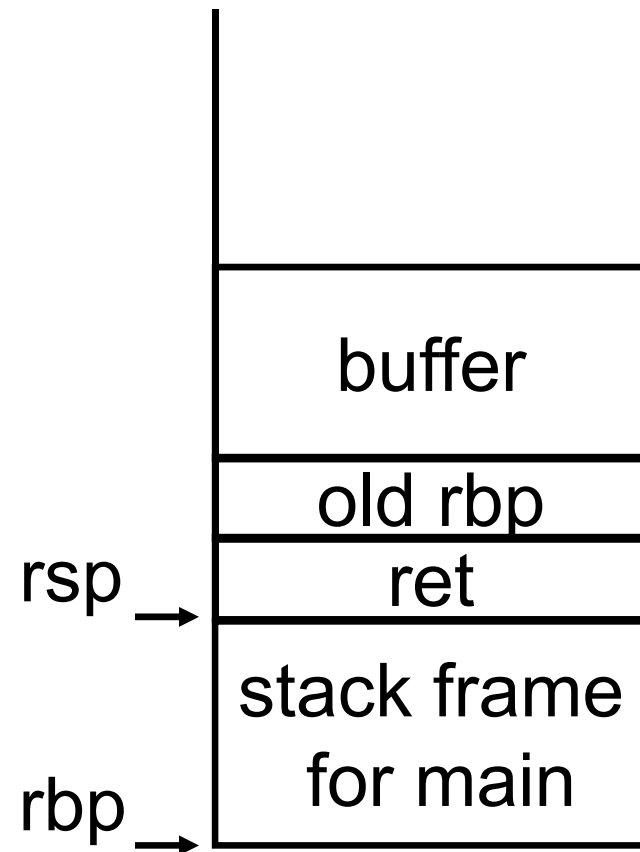


Function Return: Stacks

Before



After

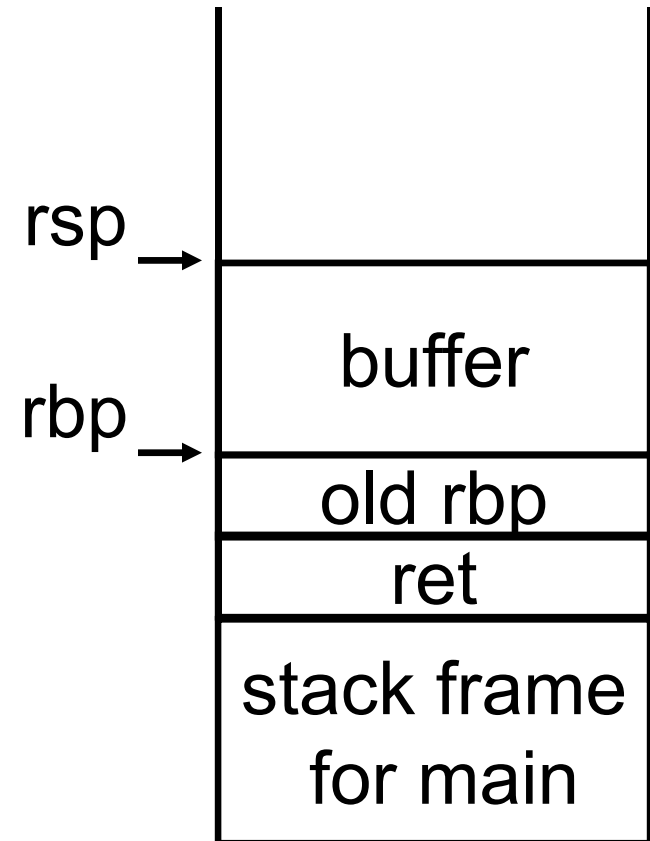




A Running Example

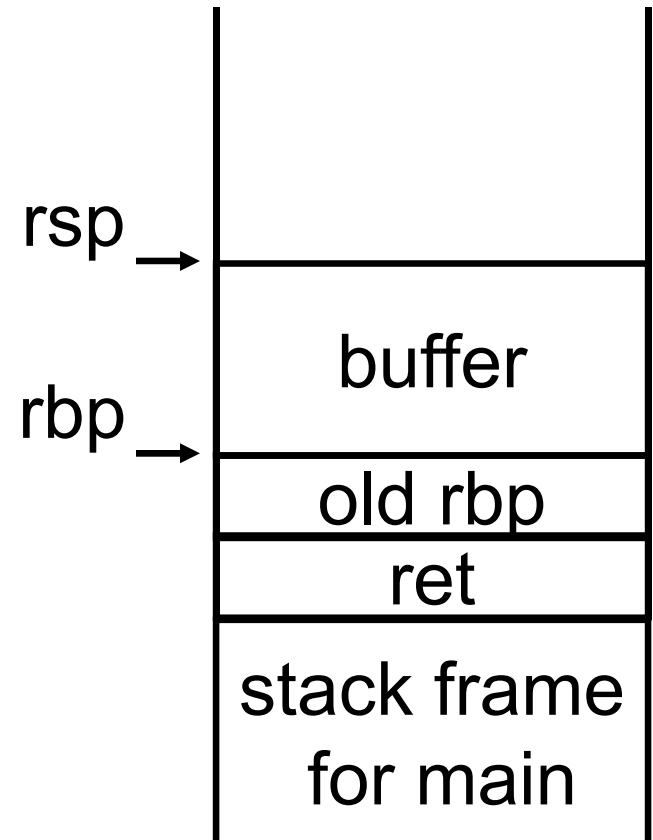
```
void function(int a, int b) {  
    char buffer[12];  
    gets(buffer);  
→   return;  
}
```

```
void main() {  
    int x;  
    x = 0;  
    function(1,2);  
    x = 1;  
    printf("%d\n",x);  
}
```



Overwriting the Return Address

```
void function(int a, int b) {  
    char buffer[12];  
    gets(buffer);  
  
    long* ret =  
        (long *) ((long)buffer+?);  
    *ret = *ret + ?;  
  
    return;  
}
```





Overwriting the Return Address

```
void function(int a, int b) {  
    char buffer[12];  
    gets(buffer);
```

```
    long* ret = (long *) ((long)buffer+40);  
    *ret = *ret + 7;
```

```
    return;  
}
```

The output will be 0

```
void main() {  
    int x;  
    x = 0;  
    function(1,2);  
    x = 1;  
    printf("%d\n",x);  
}
```

the original return address

the new return address



The Previous Attack

- Not very realistic
 - Attackers are usually not allowed to modify code
 - Threat model: the only thing they can affect is the input
 - Can they still carry out similar attacks?
 - YES, because of possible buffer overflows




Buffer Overflow

- A buffer overflow occurs when data is written outside of the boundaries of the memory allocated to a particular data structure
- Happens when buffer boundaries are neglected and unchecked
- Can be exploited to modify
 - return address on the stack
 - local variable
 - heap data structures
 - function pointer



Smashing the Stack

- Occurs when a buffer overflow overwrites data in the program stack
- Successful exploits can overwrite the return address on the stack
 - Allowing execution of arbitrary code on the targeted machine



Smashing the Stack: example.c

What happens if we input a large string?

```
./example
```

```
ffffffffffffffffffffffffffffffffffffffffffffffffffffffffffffffff
```

Segmentation fault

What Happened? The Stack is Smashed

```
void function(int a, int b) {  
    char buffer[12];  
    gets(buffer);  
    return;  
}
```

If the input is large, then
gets(buffer) will write outside
the bound of buffer, and the
return address is overwritten





Figure Out A Nasty Input

```
void function (int a, int b) {  
    char buffer[12];  
    gets(buffer);  
    return;  
}
```

```
void main() {  
    int x;  
    x = 0;  
    function(1,2);  
    x = 1;  
    printf("%d\n",x);  
}
```




A nasty input puts the return
address after x=1.

Arc injection


Injecting Code

```
void function (int a, int b) {  
    char buffer[12];  
    gets(buffer);  
    return;  
}
```

```
void main() {  
    int x;  
    x = 0;  
    function(1,2);  
    x = 1;  
    printf("%d\n",x);  
}
```




The injected code can do anything. E.g., download and install a worm



Code Injection

- Attacker creates a malicious argument—a specially crafted string that contains a pointer to malicious code provided by the attacker
- When the function returns, control is transferred to the malicious code
 - Injected code runs with the permission of the vulnerable program when the function returns.
 - Programs running with root or other elevated privileges are normally targeted
 - Programs with the setuid bit on

Injecting Shell Code



- This brings up a shell
- Attacker can execute any command in the shell
- The shell has the same privilege as the process
- Usually a process with the root privilege is attacked



Morris Worm (1988)

- Worked by exploiting known buffer-overflow vulnerabilities in sendmail, fingerd, rsh/rexec and weak passwords.
 - e.g., it exploited a gets call in fingerd
- Infected machines probe other machines for vulnerabilities
- 6000 Unix machines were infected; \$10M-\$100M cost of damage
- Robert Morris was tried and convicted of violation of Computer Fraud and Abuse Act
 - 3 years of probation; 400 hours of community service; \$10k fine
 - Had to quit his Cornell PhD

Any C(++) code acting on untrusted input is at risk

- code taking input over untrusted network
 - eg. sendmail, web browser, wireless network driver,...
- code taking input from untrusted user on multi-user system,
 - esp. services running with high privileges (as ROOT on Unix/Linux, as SYSTEM on Windows)
- code processing untrusted files
 - that have been downloaded or emailed
- also embedded software, eg. in devices with (wireless) network connection such as mobile phones with Bluetooth, wireless smartcards in new passport or OV card, airplane navigation systems, ...